

Design and Analysis of Algorithms

Lecture 08 – Strings Searching

String Searching

Content produced by persons is most efficiently stored as text.

Structured (machine-readable) text is good for interacting with automated systems.

Free text enables personal creativity, yet is open for machine manipulation.

Sound and video data is analyzed and tagged for efficient access (ie. jump to where he kisses her) and those tags are text.

Therefore finding a particular text pattern in a large text file is one of the core algorithms.

The need has been around since early 1970s, and most famous algorithms have been developed in late 70s or early 80s.

Exact matching and similar text matching are different problems but their solution approaches are generally the same.

The worst case performance of such algorithms is very slow.

String Searching

The Brute Force Approach

Represent the text file as a very long array of characters.

Practical limit is around 2 GB per text file.

Works better if your search can be limited to per line, so that your pattern does not span multiple lines.

Among this very long array of characters, find if a shorter array exists (ie. substring search).

Have a helper function to check sameness. This function compares aligned characters one by one, and.

Align the first character of the text array with the first character of the pattern array. Run the helper function.

If there is a mismatch, then shift the pattern by one.

String Searching

The Brute Force Approach

Demonstration

Pass 1 / Comparisons (until mismatch) 2

Text: **A** **D** F G A B R A C A D A B R A F G J

Pattern: **A B** R A C A D A B R A

Pass 2 / Comparisons (until mismatch) 1

Text: A **D** F G A B R A C A D A B R A F G J

Pattern: **A B** R A C A D A B R A

Pass 3 / Comparisons (until mismatch) 1

Text: A D **F** G A B R A C A D A B R A F G J

Pattern: **A B** R A C A D A B R A

String Searching

The Brute Force Approach

Let the size of the pattern be m , and the size of the text be n .

The brute force approach makes up to $(n-m+1)$ runs of the helper function. Since we assume $m \ll n$, this is roughly equal to n . This means the main algorithm is $O(n)$.

The helper function is obviously $O(m)$.

Therefore the brute force is $O(mn)$.

Most “optimized” algorithms have cases in which they degenerate into the brute force approach.

All popular algorithms have some assumptions that make the search much much faster.

String Searching

Karp-Rabin Algorithm (1987).

This is an improvement over brute force. We run the $O(m)$ helper function only when the text is similar to the pattern.

We need a specialized hash function that is **easy to compute for a shifting pattern of values** (rolling window).

An example of such a hash function is a **rolling sum**. We subtract the item that is going out of the window, and add the newly added item. This is just 2 simple operations.

Polynomial functions used in this type of scheme are called Rabin fingerprints.

Once we have such a function, we can have a first pass, calculating the hashes of all sub-strings starting with each item in the large array. This is obviously $O(n)$.

We store these hashes in a hash table for fast access.

Then we calculate the hash for the pattern. This is $O(1)$.

Once we have all hashes, we search for hash values that are similar to the hash value we are searching. We run the one-by-one comparison only on those locations.

String Searching

Karp-Rabin Algorithm (1987).

The algorithm is $O(n + km)$ where k is the number of times we need to run the one-by-one comparison. Note that the hash function may give us false positives, so we cannot safely argue that $k=1$.

The advantage of Karp-Rabin appears when we do multiple searches on the same text file over time. The hash table is re-used over and over.

For a large number of runs over time (K), we get $O((n + Kkm) / K)$ average performance. If K approaches n then this is $O((n + nkm)/n) = O(1 + km) \sim O(m)$.

This is one of the reasons it is used commonly in **plagiarism detection**.

Extract patterns in the submitted homework (many short arrays), and then search for those patterns in already hashed reference material (research papers, etc).

When there are exact matches, mark the pattern in the homework.

At the end calculate the percentage of text marked in the homework.

String Searching

Boyer-Moore Algorithm (1977).

This is one of the most efficient algorithms for generic use.

The idea is that, once a mismatch is found, we could shift more than one character. We can shift to a **position after the mismatch**. Therefore the helper function starts comparing from the end. If we get mismatches from the end, we shift much more.

There are two types of shifts introduced in the algorithm

- Bad-character shift

- Good-suffix shift

The algorithm is $O(n/m)$ in the best case, and $O(mn)$ in the worst case.

String Searching

Boyer-Moore Algorithm (1977).

Bad-character shift

The bad-character rule considers the character in the text at which the comparison process failed (assuming such a failure occurred).

The next occurrence of that character to the left in pattern is found, and a shift which brings that occurrence in line with the mismatched occurrence in text is proposed.

If the mismatched character does not occur to the left in pattern, a shift is proposed that moves the entirety of pattern past the point of mismatch.

String Searching

Boyer-Moore Algorithm (1977).

Good-suffix shift

Suppose for a given alignment of P and T , a substring t of T matches a suffix of P , but a mismatch occurs at the next comparison to the left.

Then find, if it exists, the right-most copy t' of t in P such that t' is not a suffix of P and the character to the left of t' in P differs from the character to the left of t in T . Shift P to the right so that substring t' in P aligns with substring t in T .

If t' does not exist, then shift the left end of P past the left end of t in T by the least amount so that a prefix of the shifted pattern matches a suffix of t in T .

If no such shift is possible, then shift P by n places to the right.

If an occurrence of P is found, then shift P by the least amount so that a proper prefix of the shifted P matches a suffix of the occurrence of P in T .

If no such shift is possible, then shift P by n places, that is, shift P past t .

String Searching

Quick Search (1977).

This is basically Boyer-Moore without the complex good-suffix shift.

The algorithm is $O(mn)$.

However if the alphabet is larger than the pattern (ie. 26 characters in English alphabet and a pattern of 4 letters), then the algorithm is observed to proceed very fast.